

Background Literature and Designs

Michael Buch

January 7, 2019

1 About this document

This is a collection of references and summaries to research in the field of meta-programming/supercompilation/partial evaluation

2 Examples

Examples drawn from paper on collapsing towers [1]:

- Regular expression matcher \vdash Evaluator \vdash Virtual Machine
 - Generate low-level VM code for a matcher specialized to one regex (through arbitrary number of intermediate interpreters)
- Modified evaluator \vdash Evaluator \vdash Virtual Machine
 - Modified for tracing/counting calls/be in CPS
 - Under modified semantics "interpreters become program transformers". E.g. CPS interpreter becomes CPS transformer

3 Methodologies

- Stage polymorphism [2]: "abstract over staging decisions" i.e. single program generator can produce code that is specialized in many different ways (instance of the Fourth Futamura Projection? [3])
- Multi-level base evaluator written in $\lambda \uparrow \downarrow$: supports staging operators (**polymorphic Lift**)
- Modify other interpreters: make them **stage polymorphic**, i.e. commands either evaluate code (like an interpreter) or generate code (like a translator)
- Stage only user-most interpreter: *wire tower* such that the **staging commands in L_n are interpreted directly in terms of staging commands in L_0** i.e. staging commands pass through all other layers handing down commands to layers below without performing any staging commands
- Non-reflective method: meta-circular evaluator **Pink** $=_i$ collapse arbitrary levels of "self-interpretation"
- $\lambda \uparrow \downarrow$ features:

- *run residual code*
- binding-time/stage polymorphism [4]
- preserves execution order of future-stage expressions
- does not require type system or static analysis
 - * TDPE [5] (great explanation also at [6]): **polymorphic Lift** operator turns static values into dynamic (future-stage) expressions

3.1 Towers of Interpreters Project Overview

3.1.1 Scala

- base.scala: implements definitional interpreter for $\lambda \uparrow\downarrow$

4 Results

- Able to achieve compilation of stack-machine on top the Pink evaluator (including tracing evaluator etc.)
- Compilation i.e. collapsing through explicit staging annotations requires intricate knowledge of infrastructure and does not support all data structures e.g. stacks

5 Problems

A useful analogy is the one presented in [1]: a Python interpreter running on a JavaScript emulator of a x86 CPU. What we envision (with reference to this hypothetical setting) is handling the two following cases:

1. A one-off run of a python script on top of this stack should be collapsed by bypassing the emulator interpretation
2. A continuously running emulator evaluating a continuously running python interpreter should collapse individual runs of interpretation while respecting the dynamically changing environment
 - Here a dynamically changing environment also implies side-effects that are capable of changing the semantics of interpreters within the tower at runtime
 - In literature, the closest to compiling a dynamically changing tower is [7, 1] (for a *reflective* language Black) and GraalVM [8]

To tackle the first of these problems we construct a similar yet condensed form of the setting as shown in 1.

6 Contributions

- Extension of the core $\lambda \uparrow\downarrow$ to support side effects, combining previous insights into multi-level λ [9] and work on side-effects in partial evaluators [7].
- Development a CESK-style abstract machine/ abstract interpreter for said extended $\lambda \uparrow\downarrow$




Figure 1: “Effectively functional” $\lambda \uparrow \downarrow^*$ with SECD tower above it

- Form a basis for further work on towers by providing a stage polymorphic base evaluator capable of modelling functional or imperative languages
- Mimick a practical tower through a SECD machine on top of the base evaluator and show compilation without staging commands throughout the tower
- Theoretical proposal of how one might achieve collapsing in practice

6.1 Staging

There are two types of partial evaluation methodologies [10]:

- Offline partial evaluation
- Online partial evaluation

Namin et al. [1], propose two languages Pink and Purple. Pink uses a form of online partial evaluation but requires manual staging facilities. Purple relies on LMS for automatic binding time analysis and staging which limits it to offline partial evaluation and thus relies on further optimization heuristics to achieve the same level of program specialization in the generated code as Pink.

Our language extends Pink with side effects and a stack machine that makes use of pointer like semantics for Lisp-like cons pairs. Thus we build on top of the NBE-style lift operator for staging. However, calling into the base-level lift requires knowledge about its use to be passed from the layers above. We can employ several strategies of doing this:

- A basic approach exposes the base layer staging operation to the level above. This is how the original Pink implementation works.
- At every layer *deduce* whether we need to call the underlying interpreter staging operator
 - This requires every level to include an implementation of such staging operations
- A mixture of passing staging operations to the layer below or implementing ones own operators
- Find a method of passing staging decisions through each layer in a generic way without intrusive changes to the evaluators of the layers
- Decide about calling staging operations at a particular point in the tower and apply previous points

We are interested in the last two point. In heterogeneous and practical towers a programmer does not have the liberty to introduce intrusive changes along each layer. The original Pink implementation assumes we are allowed to make arbitrary changes to evaluators. It effectively adds tags to the emitted representation of a layer above and lets the layer below infer from these tags what tag it itself should pass to the next layer, eventually calling the base-level Lift term.

6.2 Why do we want to collapse towers? (aren't they pretty?)

The main reason is performance. The key realization of partial evaluation that lead to its development is that interpreters do redundant work but we can make it so they don't. Program specialization is simple and attractive on paper but poses significant engineering challenges and has not seen widespread adoption (until recent increasingly successful work on interpreter virtualization [8]).

Binding time analysis is one of the obstacles of program specialization. The program specializer needs to decide, either automatically or with assistance from the programmer, which data to treat as static and which as dynamic. Simple divisions can lead to code explosion or inefficient code generation, or worse, to non-termination of the specializer. This problem is known as *division* and is one of the key differences between offline and online PE techniques [10].

An curious use-case for staging towers of interpreters started with the challenge of compiling the reflective language proposed by Asai et al. [7, 11]. The authors are able to compile a Lisp-like reflective language, built through the infinite tower of interpreters model [12, 13, 14], with respect to the initial semantics of the tower. Amin et al. [1] then extended this work to allow compilation of such towers under modified semantics i.e. dynamically changing behaviour of individual interpreters. An interesting consequence demonstrated in their paper is the ability to derive translators in the process of collapsing.

6.2.1 Example of Deriving Translators

A trivial but useful example is logging. Given the tower in figure ?? we want to keep the added *useful* behaviour of I_3 while removing the *unuseful* other work of interpreting an intermediate representation. The interpretation of the IR of the level above is a mere accidental consequence of design instead of a necessity. We claim this work to be *interpretative overhead* and defer its quantification by benchmarks to a later section.

Collapsing the tower achieves exactly what we wanted, base-language (here the compilation target language) expressions including logging specialized to the user-level program.

A restriction with this method is its reliance on meta-circularity and reflection and other unsafe techniques:

- we are able “inject” the logging evaluator into tower because of its meta-circularity
- we expose staging operators throughout the tower through simple string manipulation
- instrumentation relies on meta-circularity since we simply redefine how constructs are evaluated before injecting the evaluator
- modification of semantics of the tower are done via reflection (ELABORATE)

Nonetheless, it is an interesting consequence of compiling towers that we can collect side-effects in individual levels (in the original framework at no extra cost), and have an interpreter above exhibit them. Analogous to a sieve, we take a coarse-grained collection of functionality, loosen it and extract only the individual grains we want. Imagine the sieve being coated with a thin film. All grains passing through will now have the film applied to them. Here sieves are the levels below and grains are the levels above. (REFINE ANALOGY)

Our study examines this property by testing the limits of how we can get the side effects to stick to interpreters in a useful way. One could imagine optimizations, parallelization or instrumentation as possible use cases. Under certain side-effects, we may, however, reach limits in terms of security (TROJANS IN HYPERVISORS) or ability to reason about a system. We are interested in the extent of these limits.

References

- [1] N. Amin and T. Rompf, “Collapsing towers of interpreters,” *Proceedings of the ACM on Programming Languages*, vol. 2, no. POPL, p. 52, 2017.
- [2] G. Ofenbeck, T. Rompf, and M. Püschel, “Staging for generic programming in space and time,” in *ACM SIGPLAN Notices*, vol. 52, no. 12. ACM, 2017, pp. 15–28.
- [3] R. Glück, “Is there a fourth futamura projection?” in *Proceedings of the 2009 ACM SIGPLAN workshop on Partial evaluation and program manipulation*. ACM, 2009, pp. 51–60.
- [4] F. Henglein and C. Mossin, “Polymorphic binding-time analysis,” in *European Symposium on Programming*. Springer, 1994, pp. 287–301.
- [5] O. Danvy, “Type-directed partial evaluation,” in *Partial Evaluation*. Springer, 1999, pp. 367–411.
- [6] B. Grobauer and Z. Yang, “The second futamura projection for type-directed partial evaluation,” *Higher-Order and Symbolic Computation*, vol. 14, no. 2-3, pp. 173–219, 2001.
- [7] K. Asai, H. Masuhara, and A. Yonezawa, *Partial evaluation of call-by-value λ -calculus with side-effects*. ACM, 1997, vol. 32, no. 12.
- [8] T. Würthinger, C. Wimmer, A. Wöß, L. Stadler, G. Duboscq, C. Humer, G. Richards, D. Simon, and M. Wolczko, “One vm to rule them all,” in *Proceedings of the 2013 ACM international symposium on New ideas, new paradigms, and reflections on programming & software*. ACM, 2013, pp. 187–204.
- [9] F. Nielson and H. R. Nielson, “Multi-level lambda-calculi: an algebraic description,” in *Partial evaluation*. Springer, 1996, pp. 338–354.
- [10] N. D. Jones, C. K. Gomard, and P. Sestoft, *Partial evaluation and automatic program generation*. Peter Sestoft, 1993.
- [11] K. Asai, “Compiling a reflective language using metaocaml,” *ACM SIGPLAN Notices*, vol. 50, no. 3, pp. 113–122, 2015.
- [12] B. C. Smith, “Reflection and semantics in a procedural language (ph. d. thesis),” *Technical Report*, 1982.
- [13] M. Wand and D. P. Friedman, “The mystery of the tower revealed: A nonreflective description of the reflective tower,” *Lisp and Symbolic Computation*, vol. 1, no. 1, pp. 11–38, 1988.
- [14] O. Danvy and K. Malmkjaer, “Intensions and extensions in a reflective tower,” in *Proceedings of the 1988 ACM conference on LISP and functional programming*. ACM, 1988, pp. 327–341.